Mini-GPT: Ablations

16 October 2026

Contents

1 Vocabulary size . . . 1
2 Optimizer . . . 2
3 Parameterization . . . 2
4 Learning-rate schedule . . . 2
5 Batch size . . . 2
6 Depth vs width . . . 3
7 Context length . . . 3
8 Document masking . . . 3
9 QK-norm . . . 3
10 Sliding-window attention . . . 4
11 Tied vs untied embeddings . . . 4
12 MLP . . . 4
13 Init and stability . . . 4
14 Grouped-query attention . . . 5
15 Data . . . 5
16 Digit splitting . . . 5

1 Vocabulary size

2 Optimizer

3 Parameterization

4 Learning-rate schedule

5 Batch size

6 Depth vs width

7 Context length

8 Document masking

9 QK-norm

10 Sliding-window attention

11 Tied vs untied embeddings

12 MLP

13 Init and stability

14 Grouped-query attention

15 Data

16 Digit splitting